

# Assignment 1

September 24, 2023

This is a companion notebook for the book [Deep Learning with Python, Second Edition](#). For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

**If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.**

This notebook was generated for TensorFlow 2.6.

## 1 Getting started with neural networks: Classification and regression

### 1.1 Classifying movie reviews: A binary classification example

#### 1.1.1 The IMDB dataset

Loading the IMDB dataset

```
[1]: from tensorflow.keras.datasets import imdb
      (train_data, train_labels), (test_data, test_labels) = imdb.load_data(
          num_words=10000)
```

```
[2]: train_data[0]
```

```
[2]: [1,
      14,
      22,
      16,
      43,
      530,
      973,
      1622,
      1385,
      65,
      458,
      4468,
      66,
      3941,
      4,
```

173,  
36,  
256,  
5,  
25,  
100,  
43,  
838,  
112,  
50,  
670,  
2,  
9,  
35,  
480,  
284,  
5,  
150,  
4,  
172,  
112,  
167,  
2,  
336,  
385,  
39,  
4,  
172,  
4536,  
1111,  
17,  
546,  
38,  
13,  
447,  
4,  
192,  
50,  
16,  
6,  
147,  
2025,  
19,  
14,  
22,  
4,  
1920,

4613,  
469,  
4,  
22,  
71,  
87,  
12,  
16,  
43,  
530,  
38,  
76,  
15,  
13,  
1247,  
4,  
22,  
17,  
515,  
17,  
12,  
16,  
626,  
18,  
2,  
5,  
62,  
386,  
12,  
8,  
316,  
8,  
106,  
5,  
4,  
2223,  
5244,  
16,  
480,  
66,  
3785,  
33,  
4,  
130,  
12,  
16,  
38,

619,  
5,  
25,  
124,  
51,  
36,  
135,  
48,  
25,  
1415,  
33,  
6,  
22,  
12,  
215,  
28,  
77,  
52,  
5,  
14,  
407,  
16,  
82,  
2,  
8,  
4,  
107,  
117,  
5952,  
15,  
256,  
4,  
2,  
7,  
3766,  
5,  
723,  
36,  
71,  
43,  
530,  
476,  
26,  
400,  
317,  
46,  
7,

4,  
2,  
1029,  
13,  
104,  
88,  
4,  
381,  
15,  
297,  
98,  
32,  
2071,  
56,  
26,  
141,  
6,  
194,  
7486,  
18,  
4,  
226,  
22,  
21,  
134,  
476,  
26,  
480,  
5,  
144,  
30,  
5535,  
18,  
51,  
36,  
28,  
224,  
92,  
25,  
104,  
4,  
226,  
65,  
16,  
38,  
1334,  
88,

```
12,  
16,  
283,  
5,  
16,  
4472,  
113,  
103,  
32,  
15,  
16,  
5345,  
19,  
178,  
32]
```

```
[3]: train_labels[0]
```

```
[3]: 1
```

```
[4]: max([max(sequence) for sequence in train_data])
```

```
[4]: 9999
```

### Decoding reviews back to text

```
[5]: word_index = imdb.get_word_index()  
reverse_word_index = dict(  
    [(value, key) for (key, value) in word_index.items()])  
decoded_review = " ".join(  
    [reverse_word_index.get(i - 3, "?") for i in train_data[0]])
```

### 1.1.2 Preparing the data

#### Encoding the integer sequences via multi-hot encoding

```
[6]: import numpy as np  
def vectorize_sequences(sequences, dimension=10000):  
    results = np.zeros((len(sequences), dimension))  
    for i, sequence in enumerate(sequences):  
        for j in sequence:  
            results[i, j] = 1.  
    return results  
x_train = vectorize_sequences(train_data)  
x_test = vectorize_sequences(test_data)
```

```
[7]: x_train[0]
```

```
[7]: array([0., 1., 1., ..., 0., 0., 0.]
```

```
[8]: y_train = np.asarray(train_labels).astype("float32")
     y_test = np.asarray(test_labels).astype("float32")
```

### 1.1.3 Building your model

#### Model definition

```
[9]: from tensorflow import keras
     from tensorflow.keras import layers

     #Tried different combinations of layers (3,2,1) and dense units (16,32,64)

     model = keras.Sequential([
         layers.Dense(16, activation="relu"),
         layers.Dense(16, activation="relu"),
         # using tanh as activation function
         #layers.Dense(16, activation="tanh"),
         #layers.Dense(16, activation="tanh"),
         layers.Dense(1, activation="sigmoid")
     ])
```

```
[10]: ### Including regularization
```

```
[ ]: from keras import models
     from keras import layers
     from keras import regularizers

     # Tried different configuration of L1= 0.005, 0.01 and L2= 0.01, 0.02
     model = models.Sequential()
     model.add(layers.Dense(16, kernel_regularizer=regularizers.l1_l2(l1=0.001, l2=0.
         ↪02), activation='relu', input_shape=(10000,)))
     model.add(layers.Dense(16, kernel_regularizer=regularizers.l1_l2(l1=0.001, l2=0.
         ↪02), activation='relu'))
     model.add(layers.Dense(1, activation='sigmoid'))
```

```
[11]: ### Including Dropout
```

```
[12]: from keras.models import Sequential
     from keras.layers import Dense, Dropout

     # Tried different combinations of dropout values (0.1, 0.5, 0.8)
     model = Sequential()
     model.add(Dense(16, activation='relu', input_shape=(10000,)))
     model.add(Dropout(0.5))
     model.add(Dense(16, activation='relu'))
```

```
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))
```

### Compiling the model

```
[13]: model.compile(optimizer="rmsprop",
                    loss="binary_crossentropy",
                    metrics=["accuracy"])
```

```
[ ]: # Using mean_square_error as loss function
      #model.compile(optimizer="rmsprop",
      #               loss="mean_squared_error",
      #               metrics=["accuracy"])
```

### 1.1.4 Validating your approach

#### Setting aside a validation set

```
[14]: x_val = x_train[:10000]
      partial_x_train = x_train[10000:]
      y_val = y_train[:10000]
      partial_y_train = y_train[10000:]
```

#### Training your model

```
[15]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 2s 40ms/step - loss: 0.6275 - accuracy: 0.6387 - val\_loss: 0.5196 - val\_accuracy: 0.8353

Epoch 2/20

30/30 [=====] - 1s 21ms/step - loss: 0.5108 - accuracy: 0.7792 - val\_loss: 0.4247 - val\_accuracy: 0.8700

Epoch 3/20

30/30 [=====] - 1s 27ms/step - loss: 0.4304 - accuracy: 0.8312 - val\_loss: 0.3685 - val\_accuracy: 0.8692

Epoch 4/20

30/30 [=====] - 1s 21ms/step - loss: 0.3674 - accuracy: 0.8627 - val\_loss: 0.3076 - val\_accuracy: 0.8849

Epoch 5/20

30/30 [=====] - 1s 22ms/step - loss: 0.3200 - accuracy: 0.8910 - val\_loss: 0.2895 - val\_accuracy: 0.8892

Epoch 6/20

30/30 [=====] - 1s 24ms/step - loss: 0.2795 - accuracy:



```

0.9067 - val_loss: 0.2770 - val_accuracy: 0.8900
Epoch 7/20
30/30 [=====] - 1s 24ms/step - loss: 0.2444 - accuracy:
0.9237 - val_loss: 0.3006 - val_accuracy: 0.8856
Epoch 8/20
30/30 [=====] - 1s 24ms/step - loss: 0.2156 - accuracy:
0.9302 - val_loss: 0.2833 - val_accuracy: 0.8870
Epoch 9/20
30/30 [=====] - 1s 21ms/step - loss: 0.1872 - accuracy:
0.9386 - val_loss: 0.2915 - val_accuracy: 0.8891
Epoch 10/20
30/30 [=====] - 1s 21ms/step - loss: 0.1667 - accuracy:
0.9483 - val_loss: 0.3165 - val_accuracy: 0.8876
Epoch 11/20
30/30 [=====] - 1s 24ms/step - loss: 0.1486 - accuracy:
0.9549 - val_loss: 0.3496 - val_accuracy: 0.8866
Epoch 12/20
30/30 [=====] - 1s 24ms/step - loss: 0.1293 - accuracy:
0.9611 - val_loss: 0.3437 - val_accuracy: 0.8884
Epoch 13/20
30/30 [=====] - 1s 24ms/step - loss: 0.1207 - accuracy:
0.9634 - val_loss: 0.3917 - val_accuracy: 0.8877
Epoch 14/20
30/30 [=====] - 1s 24ms/step - loss: 0.1156 - accuracy:
0.9655 - val_loss: 0.3912 - val_accuracy: 0.8864
Epoch 15/20
30/30 [=====] - 1s 24ms/step - loss: 0.1016 - accuracy:
0.9691 - val_loss: 0.4151 - val_accuracy: 0.8845
Epoch 16/20
30/30 [=====] - 1s 21ms/step - loss: 0.0914 - accuracy:
0.9725 - val_loss: 0.4630 - val_accuracy: 0.8853
Epoch 17/20
30/30 [=====] - 1s 21ms/step - loss: 0.0862 - accuracy:
0.9743 - val_loss: 0.4805 - val_accuracy: 0.8867
Epoch 18/20
30/30 [=====] - 1s 20ms/step - loss: 0.0801 - accuracy:
0.9760 - val_loss: 0.5169 - val_accuracy: 0.8831
Epoch 19/20
30/30 [=====] - 1s 20ms/step - loss: 0.0753 - accuracy:
0.9758 - val_loss: 0.5220 - val_accuracy: 0.8785
Epoch 20/20
30/30 [=====] - 1s 18ms/step - loss: 0.0743 - accuracy:
0.9781 - val_loss: 0.5739 - val_accuracy: 0.8848

```

```


[16]: history_dict = history.history
      history_dict.keys()

```

```
[16]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

### Plotting the training and validation loss


```
[17]: import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



### Plotting the training and validation accuracy

```
[18]: plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
```

```
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



### Retraining a model from scratch

```
[19]: #Tried different combinations of layers (3,2,1) and dense units (16,32,64)
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    # using tanh as activation function
    #layers.Dense(16, activation="tanh"),
    #layers.Dense(16, activation="tanh"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
# Using mean_square_error as loss function
#model.compile(optimizer="rmsprop",
#              loss="mean_squared_error",
#              metrics=["accuracy"])
```

```
model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

```
Epoch 1/4
49/49 [=====] - 2s 19ms/step - loss: 0.4644 - accuracy:
0.8238
Epoch 2/4
49/49 [=====] - 1s 17ms/step - loss: 0.2671 - accuracy:
0.9063
Epoch 3/4
49/49 [=====] - 1s 14ms/step - loss: 0.2035 - accuracy:
0.9273
Epoch 4/4
49/49 [=====] - 1s 13ms/step - loss: 0.1703 - accuracy:
0.9387
782/782 [=====] - 1s 594us/step - loss: 0.2921 -
accuracy: 0.8837
```

```
[20]: results
```

```
[20]: [0.2920909821987152, 0.883679986000061]
```

### 1.1.5 Using a trained model to generate predictions on new data

```
[21]: model.predict(x_test)
```

```
[21]: array([[0.16282263],
           [0.99993783],
           [0.852559  ],
           ...,
           [0.1051057  ],
           [0.05724409],
           [0.66419065]], dtype=float32)
```